


Convolutional neural networks for visual recognition in q

(with additional q-integrated CUDA/GPU functions for
performance gains)

Ryan Sparks November 2017

Overview

- What this is all about/why I did this
- Background on visual recognition
- Rehash on neural nets
- What are convnets and important layers/general architecture of convnets
- Key techniques to improve accuracy
- How I got this working/issues encountered



<https://raweb.inria.fr/rapportsactivite/RA2007/lear/uid64.html>



q)

Overview

- Improvements I found helpful
- Setting up Google Cloud to run this (with a GPU), and hook up to a cell phone
- Using CUDA to complement q, some egs (matrix multiply, shortest path)
- Overall assessment of doing this in q vs existing languages/frameworks
- Future plans/TODOs



What's this about

- Image classification in q
- Following stanford's cs231n class, implemented using the CIFAR 10 dataset:
<https://www.cs.toronto.edu/~kriz/cifar.html> (more detail later slides)
- 10 different classes, 49,000 training images, 1000 validation images

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck

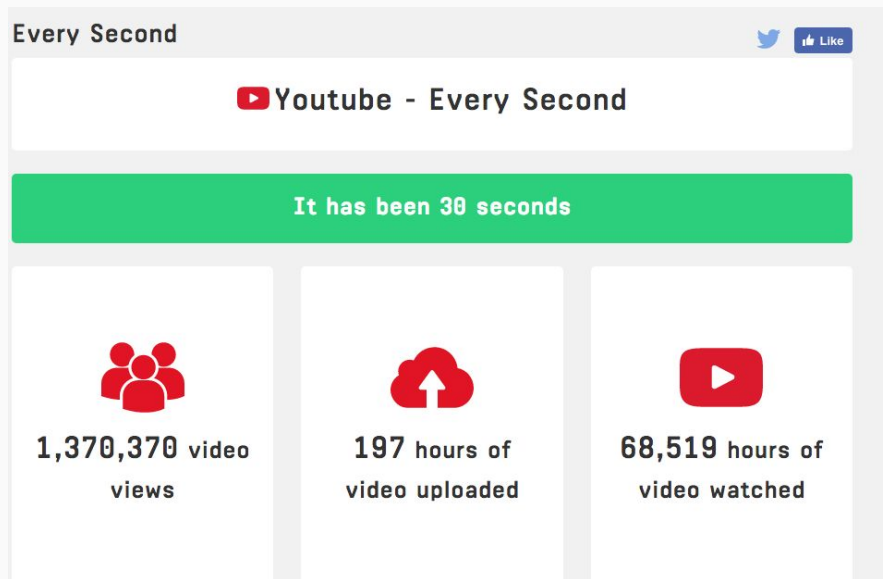


Why

- Visual recognition seems to be one of the most exciting, useful, and rapidly expanding fields in technology/AI at the moment
- Wanted to understand myself how it worked, best way for me to truly learn is write code to do something
- I like coding in q, for the same reason as everyone here (also the language I'm most proficient in)
- One of the most important/modern methods of visual recognition can be done in q - it's not just dark magic
- Leads onto other exciting areas as I'll mention later

Visual recognition - importance

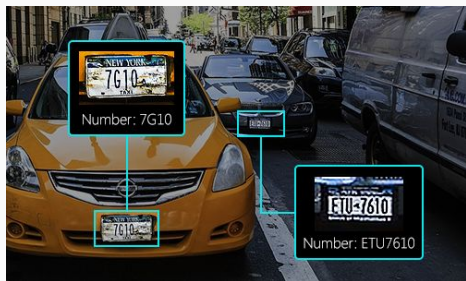
- Number of cameras in the world overtaking number of people, > 2 billion smartphones
- Up to 80% of all internet traffic is video
- Majority of data is visual data, critical, but difficult to for algos understand
- <http://www.everysecond.io/youtube> 6 hours of youtube every second
- Impossible for humans to look at and interpret all the visual data
- Convolutional neural networks recently emerged to help solve this problem




Convolutional neural networks: endless applications



<https://c.tribune.com.pk/2016/09/1191315-self-drivecar-1475231284.jpg>



<https://c.tribune.com.pk/2016/09/1191315-self-drivecar-1475231284.jpg>



<https://stanfordmlgroup.github.io/projects/ecg/>



<http://blog.twmg.com.au/wp-content/uploads/2015/09/colorized-historic-photo-15-30.jpg>




<https://research.googleblog.com/2014/09/building-deeper-understanding-of-images.html>



<https://benheubl.github.io/data%20analysis/fr/>

Rehash on neural nets

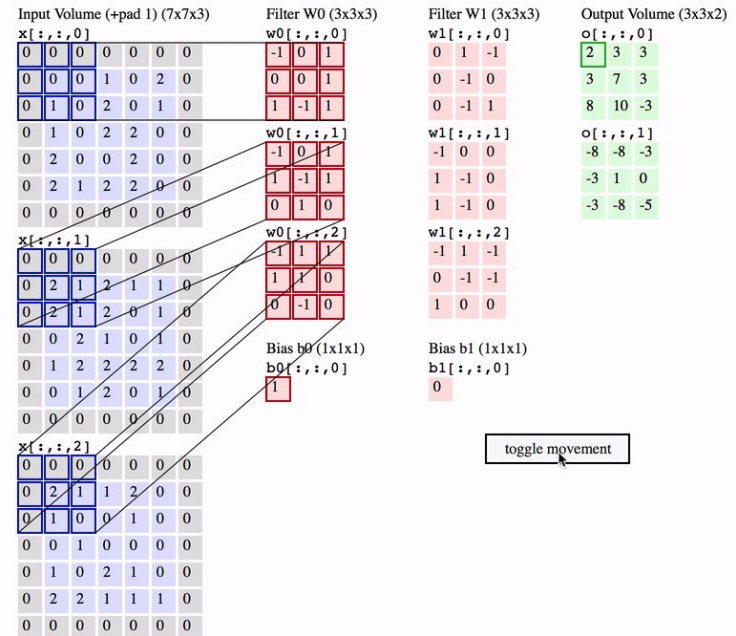
- Simplified models of brains
- Lots of neurons arranged in layers, with connections to all neurons in next layers
- Each connection has associated weight
- Training information fed through, certain neurons fire (activation function)
- Compare output to expected, then adjust weights using simple calculus, and repeat
- Primarily a matrix multiply (at least in compute time)



http://cs231n.github.io/assets/nn1/neural_net.jpeg

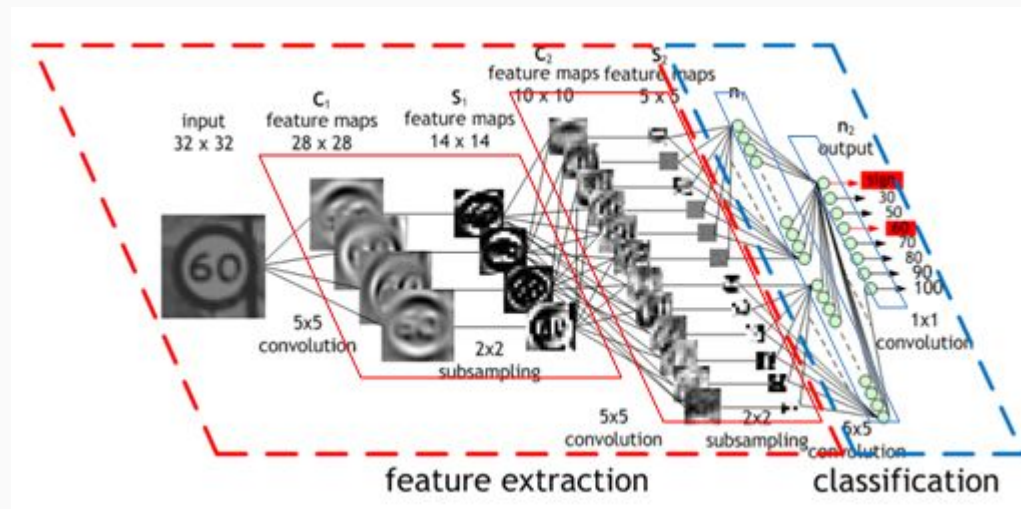
Convolutional neural networks (convnets)

- Category of neural network that is very good at image recognition
- Been around since 80's, but didn't scale well until 2012 (when GPUs and huge image sets were available e.g ImageNet)
- 4 main layers:
 - Convolution
 - Non linearity (reLu)
 - Pooling
 - Classification (fully connected)



Convnets

- Convolution step extracts features from input image
- Slides a small filter across an image, pixels are treated near each other which helps make sense of the image and reduce number of parameters
- Useful for 3D data, as it's embedded in architecture



<https://devblogs.nvidia.com/wp-content/uploads/2015/11/fig1.png>

Convnet layers

- Convolution layer: main function is equivalent to python's `np.lib.stride_tricks.as_strided`


```
q) asStrided: {[m;newshape;strides] newshape#razeo[m]@{raze x+/:raze  
y}/[reverse[strides]*til each reverse newshape]}  
q)\ts 0N! shape asStrided[50 3 34 34 #1f ; 3 3 3 50 32 32; 1156 34 1 3468  
34 1 ]  
3 3 3 50 32 32  
29 40128704
```

- Non linearity: reLu layer, easiest one (advantage over sigmoid is doesn't suffer same vanishing gradient problem when many layers): 0 |
- Pooling: form of subsampling, reduces dimensionality

```
{ [m;axes] { [x;ind] . [x;ind# (::) ;max] } / [m;axes] }
```

Convnet layers

- Fully connected: every neuron connected to every neuron in the next layer.
- So flow of information between each input dimension (pixel location) and each output class. Essentially matrix multiply: $\text{mmu} / \text{.qml} . \text{mm}$



Convnet layers

- Softmax: takes vector of arbitrary values (from fc layer) and squashes it so that they sum to 1:

```
// x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class for the ith input.
// y: list of labels, length N, where y[i] is the label for x[i] a
{[x;y]
  probs:{x%sum each x}exp x- max each x;
  loss:sum neg[log probs@'y'%count x;
  dx:@'[probs;y;-;1]%N;
  (loss;dx)
}
```

Design choices that improve accuracy

- Batch normalization: forcing gaussian distribution at the start, after conv layers, and fully connected layers (helps prevent init. problems too)
- Regularization: penalizing the squared magnitude of all parameters directly in the objective
`loss:dataLoss+0.5*d[`reg']*{x$x}razeo d@d`wParams`
- Dropout: while training, dropout is implemented by randomly setting neurons to 0, with probability p: `x*shapex#((prd[shapex:shape x]?1f)<p)%p`
- Data augmentation: Flip the training images over x-axis; Sample random crops in original image; jitter colors (better to preprocess): `@[xBatch;{neg[x div 2]?x}d`batchSize;reverse each]`




Design choices that improve accuracy

Different parameter update functions can help dramatically

- **sgd**: vanilla update, change the parameters along the negative gradient direction - $w \leftarrow w - dw * \text{config.learnRate}$
- **sgd with momentum**: similar, parameter vector will build up velocity in any direction that has consistent gradient -

```
v: (v*config.momentum) - dw*config.learnRate;  
config.config, enlist[velocity]!enlist v;  
(w+v;config)
```
- **rmsProp**: adaptive learning rate, uses moving avgs and squared gradients -

```
cache: (cache*updateDecayRate) + (1-updateDecayRate) * dx*dx;  
nextX: x - learnRate*dx%epsilon + sqrt cache;  
config[cache]: cache;  
(nextX;config)
```
- **adam**: similar/improved versions of rmsProp and adagrad, also uses momentum



Varying hyperparameters

- Initial learning rate
- Learning rate decay
- Regularization strength
- Randomization seems to be better than methodical
- Peach can be useful here


```
randomLearnRates:lrRange[0]+numRandoms?lrRange[1]-lrRange 0;  
randomRegs:regRange[0]+numRandoms?regRange[1]-regRange 0;  
(.solver.train @[d;`lr`reg;;;]@)peach randomLearnRates,'randomRegs
```

How I went about this

- Followed cs231n stanford course, all lectures on youtube, slides, github repo and notes freely available online
- Lots of googling, heaps of material online (blogs, githubs etc)
- Followed stanford's assignments, which are all in python/numpy
- Approx 50% of the code is provided the they leave you the rest to fill in
- I didn't really know python, used pyq/qpython to go back and forth between and workout what was really happening

Results

Best validation accuracy I found was slightly over 80% with a 5 layer convnet (compared to around 52% with a deep neural net). Definitely some overfitting here unfortunately:



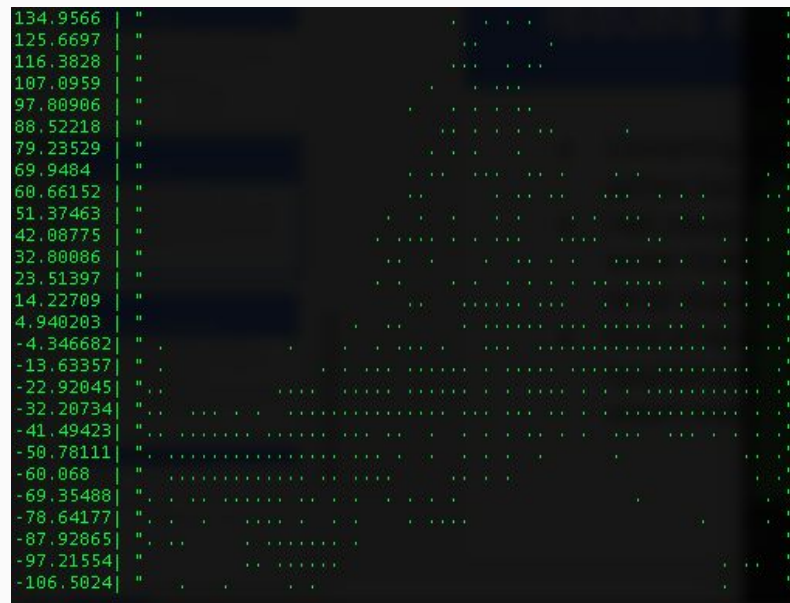
Issues encountered

- Had to learn python for this, as the course was taught in that, found pyq and qpython to be really useful
- Plenty of utils unfortunately seem to exist in numpy that aren't really well documented, had to write these from scratch in q via guess work and comparing with python.
- Wish q had multi variable assignment (for local variables inside functions, like python does), example of current work around:

```
/ convolution layer
conv_convCache:convForwardFast[x;w;b;convParam]; / returns list (conv;convCache)
/ spatial batchnorm layer
norm_normCache:spatialBatchNormForward[ conv_convCache 0;gamma;beta;bnParam]
..
cache:`convCache`normCache`reluCache!( conv_convCache 1;normCache;reluCache);
```

Issues encountered

- Converting an actual photo into an array of floats is more difficult than I had thought,
- Used python library
- 32bit w-aborting on processing the input data – so had to do some messy coding for that,
- kx fortunately gave me a temporary 64 bit license for this project
- Hard to visualize red-green-blue color photos in kdb, I tried Nick Psaris's plot function just using one color channel (as it's designed for black and white), gave up



Pictured: cat (I think?)


```

1182.797
1166.595
1156.592
1154.189
1117.986
1161.784
1085.581
1069.378
1053.176
1036.973
1028.77
1064.568
988.3649
972.1622
955.9595
939.7568
923.5541
907.3514
891.1486
874.9459
858.7432
842.5405
826.3378
810.1351
793.9324
777.7297
761.527
745.3243
729.1216
712.9189
696.7162
680.5135
664.3108
648.1081
631.9054
615.7027
599.5
583.2973
567.0946
550.8919
534.6892
518.4865
502.2838
486.0811
469.8784
453.6757
437.473
421.2703
405.0676
388.8649
372.6622
356.4595
340.2568
324.0541
307.8514
291.6486
275.4459
259.2432
243.0405
226.8378
210.6351
194.4324
178.2297
162.027
145.8243
129.6216
113.4189

```

Things I found helpful

- Deals primarily with "4 dimensional" data, i.e. lists of 3d data, eg. `2 3 4 5#til 120`.
- This was the obvious first way to do everything (can somewhat see it, in the same way you can in python, making it easy to compare)
- A big bottleneck is matrix multiply, using Andrey Zholos's qml library helped out quite a bit (e.g. 3-5 times faster than q's mmu).
- The slowest function involved a 6 level nested for loop – tried it in q, but quickly decided it needed c (ended up going from the slowest function in the whole project, to one of the fastest)

Deep loop - backward pass for convolution layer

```
// modify arg2's elements, using indexing into arg1
for(n=0;n<N;++n){
    for(c=0;c<C;++c){
        for(hh=0;hh<HH;++hh){
            for(ww=0;ww<WW;++ww){
                for(h=0;h<out_h;++h){
                    for(w=0;w<out_w;++w){
                        kF(kK(kK(kK(arg2)[n])[c])[(stride*h+hh)])[stride*w+ww] +=
                            kF(kK(kK(kK(kK(kK(arg1)[c])[hh])[ww])[n])[h])[w];
                    }
                }
            }
        }
    }
}
```

Things I found helpful

Conceptually I find it easier to work in actual matrices (and higher order list of lists), think numpy just has views on a list of data. The time to reshape adds up though!

```
q) reshapeM: { [m1;m2Shape] m2Shape#razeo m1 }
```

Eg turning a 6 dimensional matrix into a 2D one for matrix multiply.

```
q)\ts reshapeM[m;27 51200]
```

```
14 30933568
```

```
>>> m=np.random.randn(3,3,3,50,32,32)
```

```
>>> now=time.time(); res=np.reshape(m,(27,51200));time.time()-now
```

```
4.100799560546875e-05 // 0.00004 seconds
```

Things I found helpful

- Began experimenting with `gpustool` and `cuBLAS` (after Nick mentioned he'd like to try)
- Found that inputs and outputs had to be flat lists, so forced to adapt
- Was only a few weeks ago, I haven't had time to refactor the `convnet` stuff.
- Haven't had time to refactor `convnet` stuff, but have rewritten the simpler `neuralnet` stuff to use lists:

```
q) (::)m:{(x;prd[x]#1f)}3 3 3 50 32 32  
3 3 3 50 32 32  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ..  
q)reshapeMNew:@[;0;;;]  
q)reshapeMNew[m;27 51200]  
27 51200  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ..  
q)\t:10000 reshapeMNew[m;27 51200]  
9
```

Things I found helpful

- Keeping everything as a list, had to make a few additional c binaries
- E.g doing “flip” on a flat version of a matrix, without using inbuilt flip

```
q)m:{(x;prd[x]?100.)}1000 2000           // list version of matrix
(shape;list)
q)m2:(#). m                             // matrix version of m
q)\ts res:flip m2                        // normal matrix flip
54 16400672
q)\ts resSlow:raze flip (#). m           // need to avoid this
72 33176800
q)\ts resFlat:flipFlat[m 1;1000;2000]    // using c shared object flipFlat
30 16777760
q)all(resSlow~resFlat;resFlat~raze res)
1b
```

Example c func needed for list versions

```
K flipFlat(K flatm, K nrows, K ncolums){
  I i,j,rows,columns,r,c,index1,index2;
  F resvalue;
  rows=nrows->n;
  columns=ncolums->n;
  K emptyres = ktn(KF,(rows*columns));
  for(r=0;r<rows;++r){
    for(c=0;c<columns;++c){
      index1=r*columns+c;
      index2=c*rows+r;
      resvalue=kF(flatm)[index1];
      kF(emptyres)[index2]=resvalue;
    }
  };
  R(emptyres);
}
```

Things I found helpful

- Function profiling with `prof.q`
- Using some of the new inbuilt (v 3.5+) debug functions

```
/ utility, determine which parent function is calling a function
.prof.getParentFunc:`$.[;1 1 0]{1_.Q.btx@-100!`}@
```

Eg: find out which functions are calling dot, and profile them in a table, without modifying anything but the dot function:


```
dotTab:([]time:"n"$();func:`$();shape1:();shape2:())
dot:{[x;y]now:.z.n;res:.qml.mm[x;y];
    `dotTab insert (.z.n-now;.prof.getParentFunc[];enlist shape x;enlist shape y);
    res};
```

Speedup from converting neural net to use just flat lists instead of matrices

nc	timepc	timepc_old	pct	pct_old	faster
dot	4.20	189.84	67.10	81.92	1
affineBackward	1.57	17.00	7.87	1.29	1
adam	0.72	17.22	7.19	2.62	1
solver.xTrainParser	5.73	79.71	5.85	1.54	1
fullyConnectedNet.loss	4.65	11.18	5.54	0.62	1
solver.genBatch	2.53	35.30	2.54	0.54	1
solver.step	1.80	26.36	1.80	0.40	1
softmaxLoss	0.49	6.13	0.49	0.09	1
reluBackward	0.12	16.11	0.48	0.98	1
solver.checkAccuracy	8.14	8.68	0.15	0.04	1
fullyConnectedForwardPassLoop	0.02	0.02	0.14	0.01	0
affineForward	0.02	1.28	0.13	0.36	1
fullyConnectedBackwardPassLoop	0.02	0.03	0.12	0.00	1
solver.i.step	0.12	3.41	0.12	0.05	1
reluForward	0.02	1.22	0.12	0.27	1
symi	0.00	0.00	0.11	0.00	1
randArrayFlat	9.90	0.08	0.00		0
affineReluForward	0.01	0.01	0.04	0.00	1
getModelValue	0.01	0.02	0.03	0.00	1
affineReluBackward	0.01	0.40	0.02	0.02	1
..					


Using a cloud to train

- A lot of convent stuff involves letting it just sit there and "train"
- Not really practical for someone with only one laptop
- Investigated amazon cloud, was really fast but blew heaps of money
- Discovered Google cloud gave you \$300 free credit for the first year, way cheaper per hour too (but slower)
- Included my setup commands on github, as setting up a server from scratch may not be obvious
- Can take a few days to train my models



Manage cloud with cell phone

- Can setup your cell phone to ssh into it and control on the go,
- So can run q code on a pretty powerful box using your phone
- I used <https://juicessh.com/>
- <https://www.youtube.com/watch?v=mkfzyQJJoZs>
- Useful for tailing logfiles to check on training
- If you really have the urge to test out a q coding idea you can



Using GPUs for speedups



- Wanted to test out trying to implement matrix multiplication using gpus and kdb
- Since I was setting up the cloud figured I'd give it a go
- Set up a cloud with a Tesla K80 gpu (\$4k on amazon), you can still do it without paying money on google cloud, it just eats away at your \$300 credit much faster
- Have included all the cuda install steps etc. in my github

```
wget https://developer.nvidia.com/compute/cuda/9.0/Prod/local_installers/cuda-repo-ubuntu1604-9-0-local_9.0.176-1_amd64-deb
sudo apt-key add /var/cuda-repo-9-0-local/7fa2af80.pub
sudo dpkg -i cuda-repo-ubuntu1604-9-0-local_9.0.176-1_amd64-deb
sudo apt-get update
sudo apt-get -y install cuda
#check
cat /var/lib/apt/lists/*cuda*Packages | grep "Package:"
# add to bash_profile
export PATH=/usr/local/cuda-9.0/bin${PATH:+:${PATH}}
export LD_LIBRARY_PATH=/usr/local/cuda-9.0/lib64${LD_LIBRARY_PATH:+:${LD_LIBRARY_PATH}}
#install samples
cuda-install-samples-9.0.sh cudaSamples
cd cudaSamples/NVIDIA_CUDA-9.0_Samples
make
```

Compiling wasn't obvious to me:

```
nvcc --compiler-options '-fPIC -DKXVER=3 -O2' -o $QHOME/l64/gpu_mmf.so --shared -lcurand  
-lcublas gpu_mmf.cu
```

The actual code ended up being straightforward:

```
// inputs (A: list, input matrix; rA: # of rows in A; cA: # cols in A; B: input list matrix;  
rB: # rows in B; rC: # rows in B)  
extern "C" K gpu_mmf(K A, K rA, K cA, K B, K rB, K cB);  
..  
cudaMemcpy(d_A, host_memoryA, sizeA, cudaMemcpyHostToDevice);  
cudaMemcpy(d_B, host_memoryB, sizeB, cudaMemcpyHostToDevice);  
..  
// Multiply A and B on GPU  
cublasDgemm(handle, CUBLAS_OP_T, CUBLAS_OP_T, m, n, k, alpha, A, lda, B, ldb, beta, C, ldc);  
..  
// Copy the result on host memory  
cudaMemcpy(host_memoryC, d_C, nr_rows_C*nr_cols_C*sizeof(double), cudaMemcpyDeviceToHost);  
..
```

Performance was awesome!

```
q) a:1000 2000#aflat:2000000?10f
q) b:2000 3000#bflat:6000000?10f
q)\t flatres:.gpu.mm[aflat;1000;2000;bflat;2000;3000]
time to allocate host and device array mems: 1.358000ms
time to copy inputs to GPU: 9.041000ms
time to perform cublas matrix multiply: 0.024000ms
time to copy result from GPU back to host: 3.494000ms
```

40

Compared mmu and .qml.mm:

```
q)\t res2:mmu[a;b]
2051
/ using qml
q)\t res3:.qml.mm[a;b]
```

485

```
q)res2~flip 3000 1000#flatres
```

1b

Using GPUs for speedups

Had a go at the shortest paths problem from the listbox

```
__global__ void fw_kernel(const unsigned int u, const unsigned int n, int * const d){
    I v1 = blockDim.y * blockIdx.y + threadIdx.y;
    I v2 = blockDim.x * blockIdx.x + threadIdx.x;
    if (v1 < n && v2 < n){
        I newPath = d[v1 * n + u] + d[u * n + v2];
        I oldPath = d[v1 * n + v2];
        if (oldPath > newPath)
        {
            d[v1 * n + v2] = newPath; ...
        }
    }
}
```

Was approx 10x faster than the fastest q function (even with q using 6 slaves, and my lack of CUDA skills)

function	2000x2000 GPU server 0 slaves	2000x2000 on fast box+6 slaves	4000x4000 GPU server 0 slaves	4000x4000 fast box+6 slaves
bridgeq	13365 49250400	3446 32826032	134495 196802624	33971 131187376
bridgeqUDA	388 592	n/a	2890 592	n/a

Compared to other existing frameworks

- At first, even python/numpy was much faster
- After converting to use lists etc much closer
- In theory, the tools are there to be able to compete – using CUDA, customized compiled C, fast data pulling etc (kdb's forte)
- Already many highly optimized frameworks out there, built full time by teams of people at eg. google/facebook, can easily configure them for say multiple GPUs, and even TPUs.



Future work

- Some more work on recurrent neural networks
- Deep dreaming
- Image captioning
- Experiment with "minimum character-level RNN language model", there's a 110 line python script written by Andrej Karpathy (now director of AI at Tesla). Converted it to 48 lines of q, without being ridiculous/losing readability (not golf code). Reads in a text, starts "hallucinating" new text. Some of the better models have computers writing Shakespeare, baby names, writing linux source code in C, pure math theorems in Latex



<https://deepdreamgenerator.com/>



Acknowledgements/further reading

Stanford university's full course online (including lectures): <http://cs231n.stanford.edu/>

Andrej Karpathy's blog/github: <http://karpathy.github.io/>

Lee Zhen Yong's blog/github: <https://bruceoutdoors.wordpress.com/cs231n-tutorials/>

Nick Psaris's github: <https://github.com/psaris/funq>

James Neil's paper: <https://kx.com/blog/an-introduction-to-neural-networks-with-kdb/>

Thanks!

All code will be available on my
github

<https://github.com/ryantorsparks/>

rspa9428@gmail.com

